

# SLURM FOR DUMMIES



*Elaborated by:*

❖ Kevin Ender, Taype Huanca

*Email:*

❖ [kevin.taype.h@uni.pe](mailto:kevin.taype.h@uni.pe)



# “SLURM” (Simple Linux Utility for Resource Management)

Definitions:



- Slurm is an open source cluster management and job scheduling system for large and small Linux clusters
- Slurm requires no kernel modifications for its operation and is relatively self-contained.

**Slurmctld**

- Planning and administration of work.

**slurmdbd**

- Storing and recording user information, tasks, permissions.

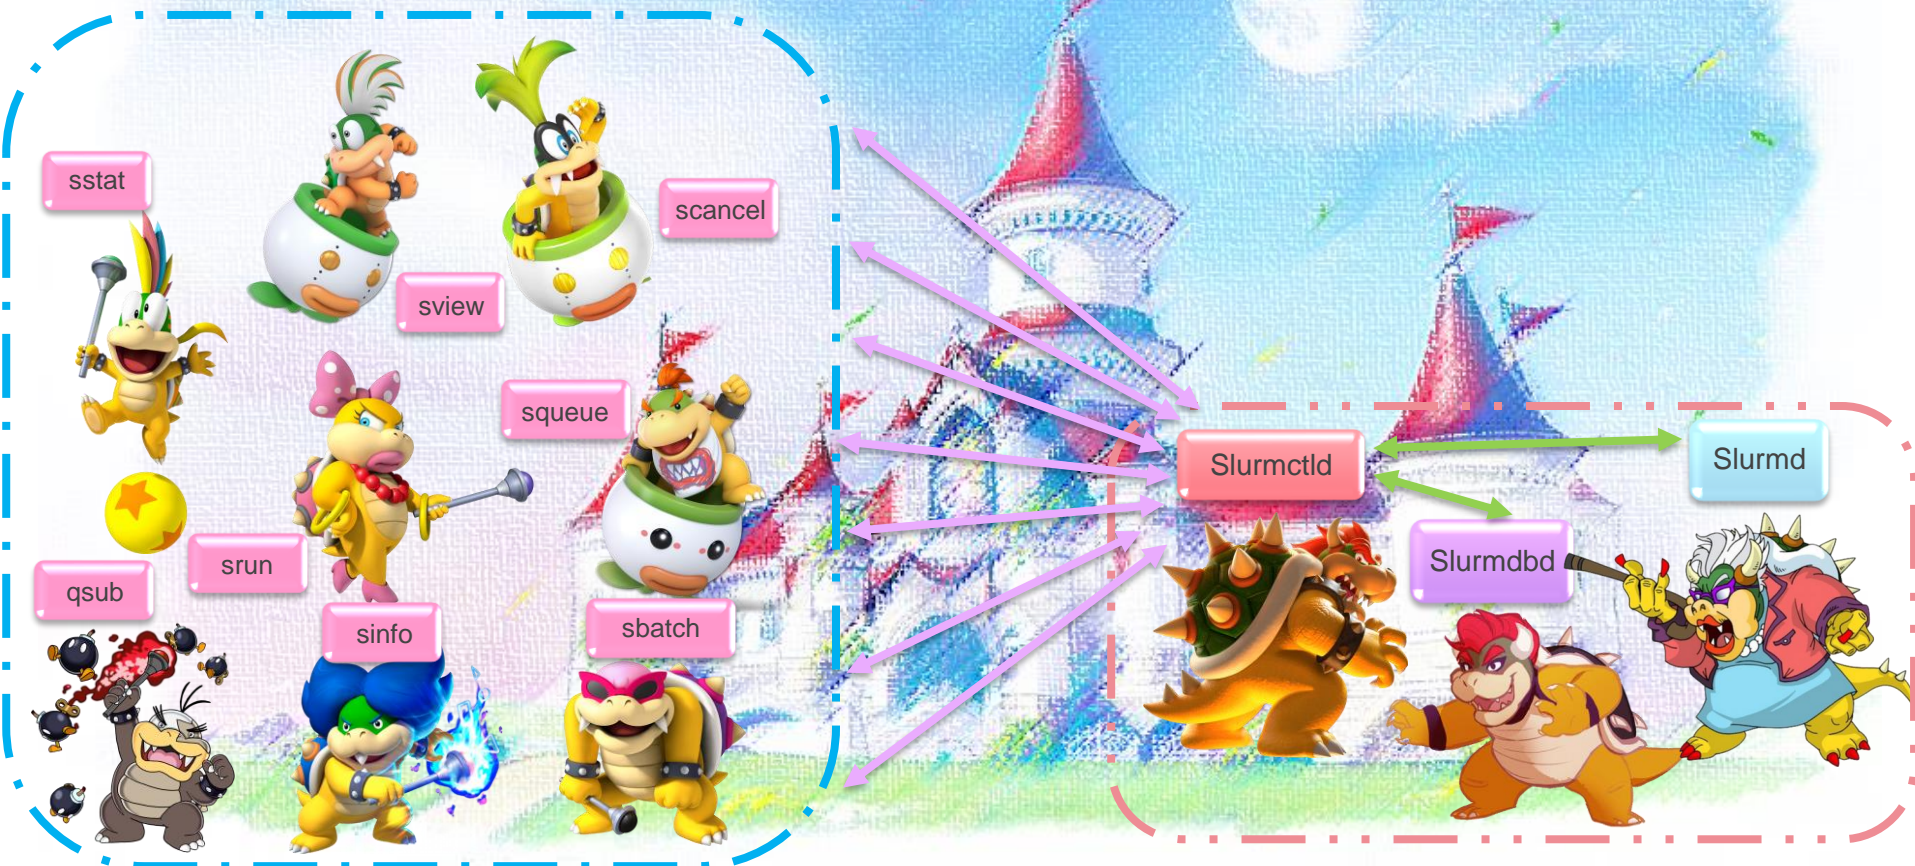
**slurmd**

- It monitors all tasks running on the compute node, accepts work (tasks), starts tasks, and kills running tasks.



**FIGURA 1.** Slurm (drink) series “Futurama”. Source: Taken from Google Images





Script:



```
#!/bin/bash
#SBATCH -J [jobname]
#SBATCH -t=[HH]:[MM]:[SS]
#SBATCH --mem=[MM] [kb/mb/gb/tb]
#SBATCH -p [partitionNAME]
#SBATCH --mail-type=[flags (NONE, BEGIN, END, FAIL, REQUEUE, ALL)]
#SBATCH --mail-user=[user]
#SBATCH -e [rutaArchivo]
#SBATCH -o [rutaArchivo]
#SBATCH -d afterok:job_id[:jobid...]
#SBATCH --export=var1=valor[,var2=valor...]
#SBATCH -N [N]
#SBATCH --tasks-per-node=[cores]
#SBATCH --gres=gpu:[numerodeGPUs]
```

...  
aplicacion y distintos pasos definidos por el usuario]



### Note

- ❖ To send the jobs to the queue manager so that it is able to know what actions it has to perform, a file is generated specifying both the resources that are required and the actions that the user wishes to execute..



